

ensemble-gjrich

April 10, 2025

1 Ensemble Machine Learning with Wine

Name: Gabriel Richards

Date: 10 Apr 2025

Introduction:

1.1 0. Imports

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import (
    RandomForestClassifier,
    AdaBoostClassifier,
    GradientBoostingClassifier,
    BaggingClassifier,
    VotingClassifier,
)
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split, cross_val_score, \
    GridSearchCV
from sklearn.metrics import (
    confusion_matrix,
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
)

# Set the state_setter variable for all instances of random_state to compare
# outcomes between seeds
```

```
# This seed influences how test_train_split cuts the data and the ensemble
↳models make their predictions
state_setter=404
```

We're going to be using pandas and numpy to perform our basic statistics and data frame handling, matplotlib and seaborn for data plotting, and sklearn for training our models.

1.2 1. Load & Inspect Data

```
[ ]: # Load the dataset
wine_data = pd.read_csv('winequality-red.csv', sep=';')

# Get information about the dataset structure
print("Dataset Information:")
wine_data.info()

# Check the shape of the dataset
print("\nDataset shape (rows, columns):")
print(wine_data.shape)
```

Our dataset has 1599 entries/rows with 12 total columns (also called features), all of which are numeric and none of which are categorical.

1.2.1 Dataset Dictionary

Here is a brief description of what each column or feature represents.

Wine Quality Features:

- **Fixed Acidity:** Primary acids in wine (mainly tartaric acid) that don't evaporate readily; contributes to structure and aging potential.
- **Volatile Acidity:** Primarily acetic acid that can evaporate; at high levels creates vinegar-like off-flavors.
- **Citric Acid:** Adds freshness and flavor to wines; found in small quantities and contributes to the wine's acidity profile.
- **Residual Sugar:** Amount of sugar remaining after fermentation stops; influences sweetness.
- **Chlorides:** Amount of salt in the wine; high levels can give a salty taste.
- **Free Sulfur Dioxide:** Unbound form of SO₂ that prevents microbial growth and oxidation.
- **Total Sulfur Dioxide:** Sum of free and bound SO₂; high levels can cause off-odors and allergic reactions.
- **Density:** Ratio of mass to volume, influenced by alcohol and sugar content; typically slightly less than water.
- **pH:** Measure of acidity on a 0-14 scale; wines typically range from 3-4 with lower values indicating acidity.
- **Sulphates:** Additive that acts as an antimicrobial and antioxidant.
- **Alcohol:** Percentage of alcohol by volume
- **Quality:** Sensory score between 0 and 10, measuring human-evaluated quality of the wine.

Let's look a little closer at what the data actually looks like.

```
[ ]: # View the first few rows
print("\nFirst 5 rows of the dataset:")
print(wine_data.head())
```

The data looks well formatted and easy to read.

```
[ ]: # Get statistical summary of numerical features
print("\nStatistical summary of features:")
print(wine_data.describe())
```

All columns have the same count, which is reassuring as we look ahead to data cleaning. Different features varied widely, with Density remaining very close to 1 but

1.3 2. Prepare the Data

Includes cleaning, feature engineering, encoding, splitting, helper functions

1.3.1 2.1 Data Cleaning & Prep

First let's clean the dataset to confirm there are no issues with it.

```
[ ]: # Check for missing values
print("\nMissing values in each column:")
print(wine_data.isnull().sum())
```

Fortunately, this data set comes clean as a whistle and has no missing, null, or visibly corrupted values.

1.3.2 2.2 Feature Engineering

2.2.1 Data Probing Let's look a little closer at the data to see how it's distributed and determine good options for our ensemble models.

```
[ ]: # Set the style for our plots
sns.set(style="whitegrid")

# Create histograms for each feature
feature_names = wine_data.columns
plt.figure(figsize=(20, 15))
for i, feature in enumerate(feature_names):
    plt.subplot(4, 3, i+1)
    sns.histplot(wine_data[feature], kde=True)
    plt.title(f'Distribution of {feature}')
    plt.xlabel(feature)
    plt.ylabel('Frequency')
plt.tight_layout()
plt.show()
```

This shows us all of the features in the data set generally follow a bell curve, although some are right skewed or left skewed.

```
[ ]: # Create boxplots for each feature
plt.figure(figsize=(20, 15))
for i, feature in enumerate(wine_data.columns[:-1]): # Exclude 'quality' for
    ↪boxplots
    plt.subplot(4, 3, i+1)
    sns.boxplot(x=wine_data[feature])
    plt.title(f'Boxplot of {feature}')
    plt.xlabel(feature)
plt.tight_layout()
plt.show()
```

These box plots (short for 'box and whisker, referring to the blue box and two verticle lines) help us grasp variance across the columns.

A boxplot visually summarizes data distribution, with the box representing the middle 50% of values and the line inside showing the median. The whiskers (horizontal lines extending from the box) reach to the smallest and largest typical values, while circles mark unusual outliers.

Boxplots make it easy to quickly compare distributions across different features, showing where most values cluster and highlighting any extreme values. When examining these wine quality plots, you can immediately see which features have wide or narrow ranges (citric acid is wide; sulphates is narrow), and which have many outliers.

Looking at them we can see that items like Citric Acid never land on a specific value across the set, whereas others like sulphates might have a strong core of typical values to establish a median (the blue box) but also have a high number of outliers (the little circles).

Let's now proceed to see how they correlate to one another. How closely is one rising or falling in value associated with each other feature?

```
[ ]: # Create a correlation matrix showing only the lower triangle
plt.figure(figsize=(12, 10))
correlation_matrix = wine_data.corr()

# Create a mask for the upper triangle
mask = np.triu(np.ones_like(correlation_matrix, dtype=bool))

# Apply the mask to the heatmap
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=0.5,
    ↪mask=mask)
plt.title('Correlation Matrix of Wine Features')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```

We can see a handful of pretty intuitive correlations, like those between citric acid and fixed. From the quality-alcohol correlation of 0.48 - quality's highest correlation - people seem to enjoy higher

alcohol-by-volume wines, with their quality generally being ranked higher.

There are some other more surprising ones - like fixed acidity's correlation with density.

There are also some interesting absences of correlation. Volatile acidity and fixed acidity have a relatively weak correlation at -0.26 for two traits which share a noun. The same is true for Sulphates and both free and total sulfur dioxide, with correlations of 0.05 and 0.04 respectively.

For our research, the bottom most row 'quality' is of greatest interest. Let's break it out and look at this a little closer.

2.2.2 Correlation with Quality

```
[ ]: quality_correlation = correlation_matrix['quality'].sort_values(ascending=False)
      print("\nFeatures sorted by correlation with quality:")
      print(quality_correlation)
```

Since we are going to be training ensemble models to predict quality, if we aim for it to be accurate we will be looking for extremities - those with very high or very low correlation.

For this reason, it would be wise to avoid including residual sugar, free sulphur dioxide, and pH in our analysis as their correlations with quality were near zero.

Volatile acidity and alcohol ranked highest with correlations of -0.39 and 0.48 respectively.

```
[ ]: # Set the style for our plots
      sns.set(style="whitegrid")
      plt.figure(figsize=(12, 10))

      # Create boxplots for features grouped by quality
      numeric_features = wine_data.columns[:-1] # All features except quality
      plt.figure(figsize=(20, 15))
      for i, feature in enumerate(numeric_features):
          plt.subplot(4, 3, i+1)
          sns.boxplot(x='quality', y=feature, data=wine_data)
          plt.title(f'{feature} by Quality Score')
      plt.tight_layout()
      plt.show()

      # Additional analysis: Feature summary grouped by quality
      quality_groups = wine_data.groupby('quality').mean()
      print("\nAverage feature values by quality score:")
      print(quality_groups)
```

2.2.3 Quality Categorization Now that we have a decent idea of how quality generally correlates with each column, we are going to categorize quality. By generalizing quality into bins, the model will have a better chance of selecting which bin - low, medium, or high quality - a wine belongs in.

```
[ ]: # Create quality categories (low, medium, high)
def quality_to_label(q):
    if q <= 4:
        return "low"
    elif q <= 6:
        return "medium"
    else:
        return "high"

# Add the quality_label column
wine_data["quality_label"] = wine_data["quality"].apply(quality_to_label)

# Check the distribution of quality categories
print("\nDistribution of quality categories:")
print(wine_data["quality_label"].value_counts())
print("\nPercentage distribution:")
print(wine_data["quality_label"].value_counts(normalize=True) * 100)

# Visualize the distribution of quality categories
plt.figure(figsize=(10, 6))
sns.countplot(x='quality_label', data=wine_data)
plt.title('Distribution of Wine Quality Categories')
plt.xlabel('Quality Category')
plt.ylabel('Count')
plt.show()

[ ]: # Create a mapping dictionary to verify consistency between the two
    ↪ categorization methods
category_mapping = {"low": 0, "medium": 1, "high": 2}

# Create numeric quality categories (0, 1, 2)
def quality_to_number(q):
    if q <= 4:
        return 0
    elif q <= 6:
        return 1
    else:
        return 2

wine_data["quality_numeric"] = wine_data["quality"].apply(quality_to_number)

# Check class distribution for the numeric categories
print("\nDistribution of numeric quality categories:")
```

```

print(wine_data["quality_numeric"].value_counts())
print("\nPercentage distribution:")
print(wine_data["quality_numeric"].value_counts(normalize=True) * 100)

# Check if the numeric categories match the label categories
wine_data["category_match"] = wine_data.apply(
    lambda row: category_mapping[row["quality_label"]] ==
    ↪row["quality_numeric"],
    axis=1
)

```

These values look right, but let's depict the numeric and label ones side by side graphically to be certain.

```

[ ]: # Display the verification results
match_percentage = (wine_data["category_match"].sum() / len(wine_data)) * 100
print(f"\nPercentage of matching categories: {match_percentage:.2f}%")
print(f"Number of mismatches: {len(wine_data) - wine_data['category_match'].
    ↪sum()}")

# If there are any mismatches, display them
if not wine_data["category_match"].all():
    print("\nMismatched entries:")
    print(wine_data[~wine_data["category_match"]][["quality", "quality_label",
    ↪"quality_numeric"]])

# Visualize the distribution of both categorization methods in one chart
plt.figure(figsize=(10, 6))

# Create a DataFrame for plotting
categories = ['Low', 'Medium', 'High']
label_counts = [
    wine_data[wine_data['quality_label'] == 'low'].shape[0],
    wine_data[wine_data['quality_label'] == 'medium'].shape[0],
    wine_data[wine_data['quality_label'] == 'high'].shape[0]
]
numeric_counts = [
    wine_data[wine_data['quality_numeric'] == 0].shape[0],
    wine_data[wine_data['quality_numeric'] == 1].shape[0],
    wine_data[wine_data['quality_numeric'] == 2].shape[0]
]

# Set up bar positions
x = np.arange(len(categories))
width = 0.35

```

```

fig, ax = plt.subplots(figsize=(12, 7))
rects1 = ax.bar(x - width/2, label_counts, width, label='Label Categories')
rects2 = ax.bar(x + width/2, numeric_counts, width, label='Numeric Categories')

# Add labels and title
ax.set_xlabel('Quality Category')
ax.set_ylabel('Count')
ax.set_title('Comparison of Quality Categorization Methods')
ax.set_xticks(x)
ax.set_xticklabels(categories)
ax.legend()

# Add value labels on bars
def add_labels(rects):
    for rect in rects:
        height = rect.get_height()
        ax.annotate('{}' .format(height),
                    xy=(rect.get_x() + rect.get_width() / 2, height),
                    xytext=(0, 3), # 3 points vertical offset
                    textcoords="offset points",
                    ha='center', va='bottom')

add_labels(rects1)
add_labels(rects2)

plt.tight_layout()
plt.show()

# Drop the verification column as it's no longer needed
wine_data.drop("category_match", axis=1, inplace=True)

```

We have a match! The label and numeric counts match and we are good to begin processing them.

1.4 3. Feature Selection and Justification

1.4.1 3.1 Analyze Feature Importance

Previously in section 2.2.2, we established Volatile acidity and alcohol ranked highest with correlations of -0.39 and 0.48 respectively. My goal is that with a comparably strong negative and strong positive correlation, the model will be able to successfully predict which of the 3 quality labels a wine belongs in.

I investigated other variables to see if there were any that had a high correlation with numeric quality, but a low correlation with the other two I intend to use.

```

[ ]: def analyze_feature_importance(data, features_list, primary_features,
    ↪target='quality'):
    """

```


Analyzes feature importance based on correlation with target and combined correlation with primary features.

```
"""
# Create a copy with only numeric columns for correlation calculation
numeric_data = data.select_dtypes(include=['number'])

# Calculate absolute correlation matrix
corr_matrix = numeric_data.corr().abs()

# Create a list to store results
results_list = []

for feature in features_list:
    # Skip if feature is already in primary features
    if feature in primary_features:
        continue

    # Get absolute correlation with quality
    quality_corr = corr_matrix.loc[feature, target]

    # Calculate combined correlation with primary features
    primary_corrs = [corr_matrix.loc[feature, pf] for pf in
    primary_features]
    combined_corr = sum(primary_corrs) / len(primary_corrs)

    # Add to results
    results_list.append({
        'Feature': feature,
        'Quality Correlation': quality_corr,
        'Primary Correlation': combined_corr
    })

# Create DataFrame from list
results = pd.DataFrame(results_list)

if len(results) == 0:
    print("No valid features found to analyze. Check feature names.")
    return pd.DataFrame(), pd.DataFrame()

# Create the two sorted tables
quality_sorted = results.sort_values('Quality Correlation',
ascending=False).reset_index(drop=True)
independence_sorted = results.sort_values('Primary Correlation',
ascending=True).reset_index(drop=True)

return quality_sorted, independence_sorted
```

```

# Define the features we want to analyze - use exact column names from your
↳dataframe
filtered_features = ['citric acid', 'total sulfur dioxide', 'density',
↳'sulphates']
primary_features = ['volatile acidity', 'alcohol']

try:
    # Calculate the tables
    quality_sorted, combined_sorted = analyze_feature_importance(wine_data,
↳filtered_features, primary_features)

    # Display the tables with formatting to 3 decimal places
    pd.set_option('display.float_format', '{:.3f}'.format)

    if not quality_sorted.empty:
        print("Features Ranked by Quality Correlation (Highest to Lowest):")
        print(quality_sorted)
        print("\nFeatures Ranked by Primary Correlation (Lowest to Highest):")
        print(combined_sorted)
except Exception as e:
    print(f"Error occurred: {e}")

    # Let's debug by examining column names
    print("\nActual column names in dataset:")
    print(wine_data.columns.tolist())

    # Check data types
    print("\nData types:")
    print(wine_data.dtypes)

```

1.4.2 3.2 Define input features and target

```

[ ]: # Define input features (X) and target (y)
X = wine_data[['volatile acidity', 'alcohol']] # Features
y = wine_data['quality_label'] # Target - using text categories instead of
↳numeric

# Display shapes to verify
print(f"X shape: {X.shape}")
print(f"y shape: {y.shape}")

# Display the first few rows of X and y to verify
print("\nFirst few rows of features (X) and target (y):")
print(X.head())
print("\nTarget classes:")
print(y.head())

```

```
# Display class distribution
print("\nClass distribution:")
print(y.value_counts())
print("\nPercentage distribution:")
print(y.value_counts(normalize=True) * 100)
```

1.5 4. Split the Data into Train and Test

1.5.1 4.1 Create train-test split

Now we will move on to split up the data for testing and training the model. We are going to stratify it to ensure the splits end up with equal low/medium/high datasets, since they are highly imbalanced (at ~4/82/14% respectively). This ensures both the test and training set maintain this distribution, which is critical for properly evaluating how well the models perform on the underrepresented low and high quality wines.

```
[ ]: # Train/test split with stratification to preserve class balance
X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,                # 80% train, 20% test
    random_state=state_setter,    # For reproducibility
    stratify=y                    # Maintain class distribution in both sets
)
```

1.5.2 4.2 Verify split characteristics

Let's confirm our split worked as intended.

```
[ ]: # Check the shapes
print(f"X_train shape: {X_train.shape}")
print(f"X_test shape: {X_test.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"y_test shape: {y_test.shape}")

# Verify stratification worked by checking class distribution
print("\nOriginal class distribution:")
print(y.value_counts(normalize=True) * 100)

print("\nTraining set class distribution:")
print(y_train.value_counts(normalize=True) * 100)

print("\nTest set class distribution:")
print(y_test.value_counts(normalize=True) * 100)
```

The numbers look good, but let's double check the actual percentages across the original set, the train set, and the test set.

```
[ ]: # Get class counts
classes = ['low', 'medium', 'high']
```

```

original_counts = [sum(wine_data['quality_label'] == cls) for cls in classes]
train_counts = [sum(y_train == cls) for cls in classes]
test_counts = [sum(y_test == cls) for cls in classes]

# Calculate percentages
total_count = len(wine_data)
train_total = len(y_train)
test_total = len(y_test)

original_percentages = [count/total_count*100 for count in original_counts]
train_percentages = [count/train_total*100 for count in train_counts]
test_percentages = [count/test_total*100 for count in test_counts]

# Calculate percentage differences
train_percentage_diffs = [train_p - orig_p for train_p, orig_p in
    ↪zip(train_percentages, original_percentages)]
test_percentage_diffs = [test_p - orig_p for test_p, orig_p in
    ↪zip(test_percentages, original_percentages)]

# Table 1: Raw Counts
counts_df = pd.DataFrame({
    'class_name': classes,
    'original_class_count': original_counts,
    'train_count': train_counts,
    'test_count': test_counts
})

# Display the tables
print("Table 1: Wine Quality Class Counts")
print(counts_df.to_string(index=False))

```

```

[ ]: # Table 2: Percentages
percentages_df = pd.DataFrame({
    'class_name': classes,
    'original_percentage': [f"{p:.2f}%" for p in original_percentages],
    'train_percentage': [f"{p:.2f}%" for p in train_percentages],
    'test_percentage': [f"{p:.2f}%" for p in test_percentages]
})

print("\nTable 2: Wine Quality Class Percentages")
print(percentages_df.to_string(index=False))

```

```

[ ]: # Table 3: Percentage Differences
differences_df = pd.DataFrame({

```

```

        'class_name': classes,
        'train_percentage_difference': [f"{p:.2f}%" for p in train_percentage_diffs],
        'test_percentage_difference': [f"{p:.2f}%" for p in test_percentage_diffs]
    })

print("\nTable 3: Percentage Differences (Split vs Original)")
print(differences_df.to_string(index=False))

```

Looks good! Our test and train distributions have put each quality class near a 0.1% difference from the original distribution. Time to move on to training.

1.6 5. Evaluate Models & Compare

Let's get set up and start feeding this data to our models.

1.6.1 5.1 Helper Function

Here we will implement a helper function to feed any of the models the data they need, and then proceed to feed it into the models. This model ensures we

```

[ ]: def evaluate_model(name, model, X_train, y_train, X_test, y_test, results):
    # Train the model
    model.fit(X_train, y_train)

    # Make predictions
    y_train_pred = model.predict(X_train)
    y_test_pred = model.predict(X_test)

    # Calculate metrics
    train_acc = accuracy_score(y_train, y_train_pred)
    test_acc = accuracy_score(y_test, y_test_pred)
    # Calculate metrics
    train_precision = precision_score(y_train, y_train_pred,
    average="weighted", zero_division=0)
    test_precision = precision_score(y_test, y_test_pred, average="weighted",
    zero_division=0)
    train_recall = recall_score(y_train, y_train_pred, average="weighted",
    zero_division=0)
    test_recall = recall_score(y_test, y_test_pred, average="weighted",
    zero_division=0)
    train_f1 = f1_score(y_train, y_train_pred, average="weighted",
    zero_division=0)
    test_f1 = f1_score(y_test, y_test_pred, average="weighted", zero_division=0)

    # Calculate gaps (to measure overfitting)
    acc_gap = train_acc - test_acc

```

```

precision_gap = train_precision - test_precision
recall_gap = train_recall - test_recall
f1_gap = train_f1 - test_f1

# Print results
print(f"\n{name} Results")
print("Confusion Matrix (Test):")
print(confusion_matrix(y_test, y_test_pred))
print(f"Train Accuracy: {train_acc:.4f}, Test Accuracy: {test_acc:.4f}, Gap:
↪ {acc_gap:.4f}")
print(f"Train Precision: {train_precision:.4f}, Test Precision:
↪ {test_precision:.4f}, Gap: {precision_gap:.4f}")
print(f"Train Recall: {train_recall:.4f}, Test Recall: {test_recall:.4f},
↪ Gap: {recall_gap:.4f}")
print(f"Train F1 Score: {train_f1:.4f}, Test F1 Score: {test_f1:.4f}, Gap:
↪ {f1_gap:.4f}")

# Store results
results.append({
    "Model": name,
    "Train Accuracy": train_acc,
    "Test Accuracy": test_acc,
    "Accuracy Gap": acc_gap,
    "Train Precision": train_precision,
    "Test Precision": test_precision,
    "Precision Gap": precision_gap,
    "Train Recall": train_recall,
    "Test Recall": test_recall,
    "Recall Gap": recall_gap,
    "Train F1": train_f1,
    "Test F1": test_f1,
    "F1 Gap": f1_gap
})

```

1.6.2 5.2 Train models

At long last! Let's proceed to train the models and look at their basic results.

```

[ ]: # Initialize results list
results = []

# 1. Random Forest
evaluate_model(
    "Random Forest (100)",
    RandomForestClassifier(n_estimators=100, random_state=state_setter),
    X_train,
    y_train,

```

```

        X_test,
        y_test,
        results
    )

# 2. Random Forest (200, max_depth=10)
evaluate_model(
    "Random Forest (200, max_depth=10)",
    RandomForestClassifier(n_estimators=200, max_depth=10,
        ↪random_state=state_setter),
    X_train,
    y_train,
    X_test,
    y_test,
    results
)

# 3. AdaBoost
evaluate_model(
    "AdaBoost (100)",
    AdaBoostClassifier(n_estimators=100, random_state=state_setter),
    X_train,
    y_train,
    X_test,
    y_test,
    results
)

```

```

[ ]: # 4. AdaBoost (200, lr=0.5)
evaluate_model(
    "AdaBoost (200, lr=0.5)",
    AdaBoostClassifier(n_estimators=200, learning_rate=0.5,
        ↪random_state=state_setter),
    X_train,
    y_train,
    X_test,
    y_test,
    results
)

# 5. Gradient Boosting
evaluate_model(
    "Gradient Boosting (100)",
    GradientBoostingClassifier(
        n_estimators=100, learning_rate=0.1, max_depth=3,
        ↪random_state=state_setter
    ),

```

```

X_train,
y_train,
X_test,
y_test,
results
)

# 6. Voting Classifier (DT, SVM, NN)
voting1 = VotingClassifier(
    estimators=[
        ("DT", DecisionTreeClassifier(random_state=state_setter)),
        ("SVM", SVC(probability=True, random_state=state_setter)),
        ("NN", MLPClassifier(hidden_layer_sizes=(50,), max_iter=1000,
↪random_state=state_setter)),
    ],
    voting="soft"
)
evaluate_model(
    "Voting (DT + SVM + NN)",
    voting1,
    X_train,
    y_train,
    X_test,
    y_test,
    results
)

```

```

[ ]: # 7. Voting Classifier (RF, LR, KNN)
voting2 = VotingClassifier(
    estimators=[
        ("RF", RandomForestClassifier(n_estimators=100,
↪random_state=state_setter)),
        ("LR", LogisticRegression(max_iter=1000, random_state=state_setter)),
        ("KNN", KNeighborsClassifier(n_neighbors=5)),
    ],
    voting="soft"
)
evaluate_model(
    "Voting (RF + LR + KNN)",
    voting2,
    X_train,
    y_train,
    X_test,
    y_test,
    results
)

```



```

# 8. Bagging
evaluate_model(
    "Bagging (DT, 100)",
    BaggingClassifier(
        estimator=DecisionTreeClassifier(random_state=state_setter),
        n_estimators=100,
        random_state=state_setter
    ),
    X_train,
    y_train,
    X_test,
    y_test,
    results
)

# 9. MLP Classifier
evaluate_model(
    "MLP Classifier",
    MLPClassifier(hidden_layer_sizes=(100,), max_iter=1000,
        random_state=state_setter),
    X_train,
    y_train,
    X_test,
    y_test,
    results
)

```

```

[ ]: # Convert results to DataFrame and sort by test accuracy
results_df = pd.DataFrame(results)
results_df = results_df.sort_values("Test Accuracy", ascending=False)
print("\nSummary of All Models (Sorted by Test Accuracy):")
display(results_df)

```

```

[ ]: # Set up consistent colors for each model to use across all charts
model_colors = {
    "Voting (DT + SVM + NN)": "#1f77b4",
    "Random Forest (200, max_depth=10)": "#ff7f0e",
    "Random Forest (100)": "#2ca02c",
    "MLP Classifier": "#d62728",
    "Voting (RF + LR + KNN)": "#9467bd",
    "Bagging (DT, 100)": "#8c564b",
    "AdaBoost (200, lr=0.5)": "#e377c2",
    "Gradient Boosting (100)": "#7f7f7f",
    "AdaBoost (100)": "#bcbd22"
}

# Sort models by test accuracy for consistent ordering

```

```

sorted_results = results_df.sort_values('Test Accuracy', ascending=False)
model_names = sorted_results['Model']
x = np.arange(len(model_names))
width = 0.25 # Slightly wider bars

# Create 4 separate figures for the different metric groups
plt.figure(figsize=(14, 8))
plt.subplot(2, 2, 1)
# Accuracy metrics
for i, model in enumerate(model_names):
    row = sorted_results[sorted_results['Model'] == model].iloc[0]
    plt.bar(i-width, row['Train Accuracy'], width, color=model_colors[model],
    ↪alpha=0.7, label=model if i == 0 else "")
    plt.bar(i, row['Test Accuracy'], width, color=model_colors[model], label="")
    plt.bar(i+width, row['Accuracy Gap'], width, color=model_colors[model],
    ↪alpha=0.4, label="")

plt.ylabel('Accuracy')
plt.title('Accuracy Metrics')
plt.xticks([])
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.legend(['Train', 'Test', 'Gap'], loc='upper right')

plt.subplot(2, 2, 2)
# Precision metrics
for i, model in enumerate(model_names):
    row = sorted_results[sorted_results['Model'] == model].iloc[0]
    plt.bar(i-width, row['Train Precision'], width, color=model_colors[model],
    ↪alpha=0.7)
    plt.bar(i, row['Test Precision'], width, color=model_colors[model])
    plt.bar(i+width, row['Precision Gap'], width, color=model_colors[model],
    ↪alpha=0.4)

plt.ylabel('Precision')
plt.title('Precision Metrics')
plt.xticks([])
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.legend(['Train', 'Test', 'Gap'], loc='upper right')

plt.subplot(2, 2, 3)
# Recall metrics
for i, model in enumerate(model_names):
    row = sorted_results[sorted_results['Model'] == model].iloc[0]
    plt.bar(i-width, row['Train Recall'], width, color=model_colors[model],
    ↪alpha=0.7)
    plt.bar(i, row['Test Recall'], width, color=model_colors[model])

```

```

plt.bar(i+width, row['Recall Gap'], width, color=model_colors[model],
        alpha=0.4)

plt.ylabel('Recall')
plt.title('Recall Metrics')
plt.xticks(range(len(model_names)), model_names, rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.legend(['Train', 'Test', 'Gap'], loc='upper right')

plt.subplot(2, 2, 4)
# F1 metrics
for i, model in enumerate(model_names):
    row = sorted_results[sorted_results['Model'] == model].iloc[0]
    plt.bar(i-width, row['Train F1'], width, color=model_colors[model], alpha=0.
        7)
    plt.bar(i, row['Test F1'], width, color=model_colors[model])
    plt.bar(i+width, row['F1 Gap'], width, color=model_colors[model], alpha=0.4)

plt.ylabel('F1 Score')
plt.title('F1 Score Metrics')
plt.xticks(range(len(model_names)), model_names, rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.legend(['Train', 'Test', 'Gap'], loc='upper right')

plt.tight_layout()
plt.show()

```

```

[ ]: # Create a heatmap for easier visual comparison
plt.figure(figsize=(12, 8))
# Prepare data for heatmap
heatmap_data = sorted_results[['Model', 'Test Accuracy', 'Test Precision',
    'Test Recall', 'Test F1']]
heatmap_data = heatmap_data.set_index('Model')
sns.heatmap(heatmap_data, annot=True, cmap='YlGnBu', fmt='.3f', linewidths=.5)
plt.title('Test Metrics Across Models')
plt.tight_layout()
plt.show()

```

```

[ ]: # Create improved scatter plots with better readability
fig, axes = plt.subplots(2, 2, figsize=(16, 12))
metrics = [
    ('Accuracy', 'Train Accuracy', 'Test Accuracy'),
    ('Precision', 'Train Precision', 'Test Precision'),
    ('Recall', 'Train Recall', 'Test Recall'),
    ('F1 Score', 'Train F1', 'Test F1')
]

```

```

# Use model colors for consistency
for i, (title, train_metric, test_metric) in enumerate(metrics):
    ax = axes[i//2, i%2]

    # Plot the diagonal line (perfect agreement between train and test)
    min_val = min(results_df[train_metric].min(), results_df[test_metric].
    ↪min()) - 0.05
    max_val = max(results_df[train_metric].max(), results_df[test_metric].
    ↪max()) + 0.05
    ax.plot([min_val, max_val], [min_val, max_val], '--', color='gray')

    # Create a mapping for abbreviated model names
    model_abbrevs = {
        "Voting (DT + SVM + NN)": "Voting (DT+SVM+NN)",
        "Random Forest (200, max_depth=10)": "RF (200, depth=10)",
        "Random Forest (100)": "RF (100)",
        "MLP Classifier": "MLP",
        "Voting (RF + LR + KNN)": "Voting (RF+LR+KNN)",
        "Bagging (DT, 100)": "Bagging",
        "AdaBoost (200, lr=0.5)": "AdaBoost (200)",
        "Gradient Boosting (100)": "Gradient Boost",
        "AdaBoost (100)": "AdaBoost (100)"
    }

    # Plot each model as a point with consistent colors
    for j, (_, row) in enumerate(sorted_results.iterrows()):
        model = row['Model']
        ax.scatter(row[train_metric], row[test_metric], ↵
        ↪color=model_colors[model], s=100, label=model if i == 0 else "")

        # Add text labels with better visibility
        ax.annotate(
            model_abbrevs[model],
            (row[train_metric], row[test_metric]),
            fontsize=10,                                # Medium font size
            xytext=(7, 0),                               # Offset text slightly
            textcoords='offset points',
            weight='bold',                                # Make text bold
            color='black',                                # Text color
            bbox=dict(boxstyle="round,pad=0.3", fc=model_colors[model], alpha=0.
            ↪3) # Background for text
        )

    ax.set_title(f'Train vs Test {title}', fontsize=14)
    ax.set_xlabel(f'Train {title}', fontsize=12)
    ax.set_ylabel(f'Test {title}', fontsize=12)
    ax.grid(True)

```

```

ax.tick_params(labelsize=10)

plt.tight_layout()
plt.show()

```

```

[ ]: # Grid search to find optimal hyperparameters
param_grid = {
    'n_estimators': [50, 100, 150],
    'max_depth': [None, 5, 10],
    'min_samples_split': [2, 5, 10]
}

grid_search = GridSearchCV(
    RandomForestClassifier(random_state=state_setter),
    param_grid=param_grid,
    cv=5,                                     # 5-fold cross-validation
    scoring='f1_weighted',                   # Optimize for F1 score
    n_jobs=-1                                # Use all available cores
)

grid_search.fit(X_train, y_train)
best_model = grid_search.best_estimator_

# Evaluate the tuned model
evaluate_model(
    "RF (Cross-Validated)",
    best_model,
    X_train, y_train, X_test, y_test, results
)

print(f"Best hyperparameters: {grid_search.best_params_}")

```

1.7 6. Compare Results

1.8 7. Conclusions and Insights